



人人都爱 DataFrame

Pandas 到 Mars 的进阶之路

何开圣



目录 CONTENTS



Python 数据科学生态

Pandas under the hood

更高效的 Pandas


更大规模的数据分析



1 Python 数据科学生态

Google Trends



PyConChina
2019





pandas

$$y_{it} = \beta' x_{it} + \mu_i + \epsilon_{it}$$



NumPy





Pythonic

完善的 API

数据分析

DataFrame

开源

强大的社区

数据清洗



什么是 Pandas

Series

- 一维数组数据
- 一个与数组关联的数据标签（索引）

```
In [12]: s = pd.Series([1, 3, 5, np.nan, 6, 8], name='col')
```

```
In [13]: s
```

```
Out[13]:
```

```
0    1.0
```

```
1    3.0
```

```
2    5.0
```

```
3    NaN
```

```
4    6.0
```

```
5    8.0
```

```
Name: col, dtype: float64
```

DataFrame

- 表格数据
- 二维有标签的数据
- 看作是一个Series的字典（所有Series共享一个索引）

```
In [8]: df
```

```
Out[8]:
```

	A	B	C	D
2013-01-01	0.469112	-0.282863	-1.509059	-1.135632
2013-01-02	1.212112	-0.173215	0.119209	-1.044236
2013-01-03	-0.861849	-2.104569	-0.494929	1.071804
2013-01-04	0.721555	-0.706771	-1.039575	0.271860
2013-01-05	-0.424972	0.567020	0.276232	-1.087401
2013-01-06	-0.673690	0.113648	-1.478427	0.524988





IO tools

- CSV & text files
- JSON
- HTML
- Excel files
- OpenDocument Spreadsheets
-

Indexing and selecting data

- Attribute access
- Slicing ranges
- Selection by label
- Selection by position
- Selection by callable
-

Group By: split-apply-combine

- Splitting an object into groups
- Iterating through groups
- Selecting a group
- Aggregation
- Transformation
-

Merge, join, and concatenate

- Concatenating objects
- Database-style DataFrame or named Series joining/merging
- Timeseries friendly merging

Computational tools

- Statistical functions
- Window Functions
- Aggregation
- Expanding windows
- Exponentially weighted windows

- **Visualization**
- **Working with text data**
- **Working with missing data**
- **Sparse data structure**





创建对象

```
In [1]: import pandas as pd
```

```
In [2]: import numpy as np
```

```
In [3]: dates = pd.date_range('20191019', periods=6)
```

```
In [4]: df = pd.DataFrame(np.random.rand(6, 4), index=dates,  
columns=list('abcd'))
```

```
In [5]: df
```

```
Out[5]:
```

	a	b	c	d
2019-10-19	0.343870	0.597608	0.873940	0.419012
2019-10-20	0.760962	0.296669	0.315243	0.836361
2019-10-21	0.294613	0.652080	0.534011	0.484261
2019-10-22	0.728239	0.928620	0.595617	0.557358
2019-10-23	0.031837	0.029126	0.982403	0.459137
2019-10-24	0.290760	0.936916	0.416385	0.646032





查看数据

```
In [8]: df.head(3)
```

```
Out[8]:
```

	a	b	c	d
2019-10-19	0.343870	0.597608	0.873940	0.419012
2019-10-20	0.760962	0.296669	0.315243	0.836361
2019-10-21	0.294613	0.652080	0.534011	0.484261

```
In [9]: df.tail(3)
```

```
Out[9]:
```

	a	b	c	d
2019-10-22	0.728239	0.928620	0.595617	0.557358
2019-10-23	0.031837	0.029126	0.982403	0.459137
2019-10-24	0.290760	0.936916	0.416385	0.646032

```
In [10]: df.index
```

```
Out[10]:  
DatetimeIndex(['2019-10-19',  
               '2019-10-20', '2019-10-21',  
               '2019-10-22',  
               '2019-10-23',  
               '2019-10-24'],  
              dtype='datetime64[ns]', freq='D')
```

```
In [11]: df.columns
```

```
Out[11]: Index(['a', 'b', 'c',  
               'd'], dtype='object')
```

```
In [12]: df.dtypes
```

```
Out[12]:  
a    float64  
b    float64  
c    float64  
d    float64  
dtype: object
```





查看数据

```
In [13]: df.describe()
```

```
Out[13]:
```

	a	b	c	d
count	6.000000	6.000000	6.000000	6.000000
mean	0.408380	0.573503	0.619600	0.567027
std	0.282612	0.357188	0.260056	0.154451
min	0.031837	0.029126	0.315243	0.419012
25%	0.291723	0.371904	0.445792	0.465418
50%	0.319242	0.624844	0.564814	0.520810
75%	0.632147	0.859485	0.804359	0.623863
max	0.760962	0.936916	0.982403	0.836361

```
In [14]: df.sort_values(by='c')
```

```
Out[14]:
```

	a	b	c	d
2019-10-20	0.760962	0.296669	0.315243	0.836361
2019-10-24	0.290760	0.936916	0.416385	0.646032
2019-10-21	0.294613	0.652080	0.534011	0.484261
2019-10-22	0.728239	0.928620	0.595617	0.557358
2019-10-19	0.343870	0.597608	0.873940	0.419012
2019-10-23	0.031837	0.029126	0.982403	0.459137





选择数据

```
In [17]: df.loc['2019-10-20':'2019-10-22', ['b', 'c']]
```

```
Out[17]:
```

	b	c
2019-10-20	0.296669	0.315243
2019-10-21	0.652080	0.534011
2019-10-22	0.928620	0.595617

```
In [15]: df['a']
```

```
Out[15]:
```

2019-10-19	0.343870
2019-10-20	0.760962
2019-10-21	0.294613
2019-10-22	0.728239
2019-10-23	0.031837
2019-10-24	0.290760

Freq: D, Name: a, dtype: float64

```
In [16]: df[:2]
```

```
Out[16]:
```

	a	b	c	d
2019-10-19	0.343870	0.597608	0.873940	0.419012
2019-10-20	0.760962	0.296669	0.315243	0.836361

```
In [18]: df.iloc[3:5,0:2]
```

```
Out[18]:
```

	a	b
2019-10-22	0.728239	0.928620
2019-10-23	0.031837	0.029126





选择数据

```
In [19]: df[df['a'] > 0.5]
```

```
Out[19]:
```

	a	b	c	d
2019-10-20	0.760962	0.296669	0.315243	0.836361
2019-10-22	0.728239	0.928620	0.595617	0.557358

```
In [20]: df[df < 0.1]
```

```
Out[20]:
```

	a	b	c	d
2019-10-19	NaN	NaN	NaN	NaN
2019-10-20	NaN	NaN	NaN	NaN
2019-10-21	NaN	NaN	NaN	NaN
2019-10-22	NaN	NaN	NaN	NaN
2019-10-23	0.031837	0.029126	NaN	NaN
2019-10-24	NaN	NaN	NaN	NaN





走进 Pandas

一些计算

```
In [21]: df.mean()
```

```
Out[21]:
```

```
a      0.408380
```

```
b      0.573503
```

```
c      0.619600
```

```
d      0.567027
```

```
dtype: float64
```

```
In [22]: df.mean(axis=1)
```

```
Out[22]:
```

```
2019-10-19      0.558608
```

```
2019-10-20      0.552309
```

```
2019-10-21      0.491241
```

```
2019-10-22      0.702459
```

```
2019-10-23      0.375626
```

```
2019-10-24      0.572523
```

```
Freq: D, dtype: float64
```

```
In [29]: s = pd.Series([1, 2, 3],  
index=df.index[2:5])
```

```
In [30]: s
```

```
Out[30]:
```

```
2019-10-21      1
```

```
2019-10-22      2
```

```
2019-10-23      3
```

```
Freq: D, dtype: int64
```

```
In [31]: df.add(s, axis='index')
```

```
Out[31]:
```

	a	b	c	d
2019-10-19	NaN	NaN	NaN	NaN
2019-10-20	NaN	NaN	NaN	NaN
2019-10-21	1.294613	1.652080	1.534011	1.484261
2019-10-22	2.728239	2.928620	2.595617	2.557358
2019-10-23	3.031837	3.029126	3.982403	3.459137
2019-10-24	NaN	NaN	NaN	NaN



走进 Pandas

合并 (SQL join)

```
In [32]: left = pd.DataFrame({'key':  
    ['foo', 'foo'], 'lval': [1, 2]})
```

```
In [33]: right = pd.DataFrame({'key':  
    ['foo', 'foo'], 'rval': [4, 5]})
```

```
In [34]: left
```

```
Out[34]:
```

	key	lval
0	foo	1
1	foo	2

```
In [35]: right
```

```
Out[35]:
```

	key	rval
0	foo	4
1	foo	5

```
In [36]: pd.merge(left, right, on='key')
```

```
Out[36]:
```

	key	lval	rval
0	foo	1	4
1	foo	1	5
2	foo	2	4
3	foo	2	5

```
In [37]: left = pd.DataFrame({'key':  
    ['foo', 'bar'], 'lval': [1, 2]})
```

```
In [38]: right = pd.DataFrame({'key':  
    ['foo', 'bar'], 'rval': [4, 5]})
```

```
In [39]: left
```

```
Out[39]:
```

	key	lval
0	foo	1
1	bar	2

```
In [40]: right
```

```
Out[40]:
```

	key	rval
0	foo	4
1	bar	5

```
In [41]: pd.merge(left, right, on='key')
```

```
Out[41]:
```

	key	lval	rval
0	foo	1	4
1	bar	2	5





走进 Pandas

分组 (groupby)

```
In [42]: df
```

```
Out[42]:
```

	A	B	C	D
0	foo	one	-0.449026	0.642683
1	bar	one	-0.074552	-0.410214
2	foo	two	-0.102928	-1.136713
3	bar	three	0.102498	0.836924
4	foo	two	-0.543563	-0.322754
5	bar	two	0.648193	-0.510235
6	foo	one	0.311538	-0.064786
7	foo	three	-1.135822	-1.399072

```
In [43]: df.groupby('A').sum()
```

```
Out[43]:
```

	C	D
A		
bar	0.676139	-0.083525
foo	-1.919801	-2.689642

- 分组
- 每个分组应用某个 **function**
- 融合最后的结果



Pandas 周边生态 —— Pandas profiling

`describe()` 实现了基础的统计信息， `pandas-profiling` 是数据分析的神器，满足了大部分的统计需求

不再需要纠结使用什么函数，分析哪些数据，直接在网站展示

```
import pandas as pd
import pandas_profiling
data = pd.read_csv("")
profile = pandas_profiling.ProfileReport(data)
profile.to_file('report.html')
```

Pandas profiling

Overview: 数据量，数据类型，warning（缺失值以及零值）

Dataset info

Number of variables	14
Number of observations	45726
Missing cells	29703 (4.6%)
Duplicate rows	0 (0.0%)
Total size in memory	4.6 MiB
Average record size in memory	105.0 B

Variables types

Numeric	4
Categorical	5
Boolean	1
Date	1
URL	0
Text (Unique)	1
Rejected	2
Unsupported	0

Warnings

GeoLocation	has a high cardinality: 17101 distinct values
GeoLocation	has 7315 (16.0%) missing values
mass_(g)	is highly skewed ($\gamma_1 = 76.91847245$)
recclass	has a high cardinality: 466 distinct values
reclat	has 6438 (14.1%) zeros
reclat	has 7315 (16.0%) missing values
reclat_city	is highly correlated with reclat ($\rho = 0.9942518712$)
reclong	has 6214 (13.6%) zeros
reclong	has 7315 (16.0%) missing values
source	has constant value "NASA"

Warning
Missing
Skewed
Warning
Zeros
Missing
Rejected
Zeros
Missing
Rejected

Pandas profiling

Variables: 每个变量的统计信息

recclass

Categorical

Distinct count

Unique (%)

Missing (%)

Missing (n)

466

1.0%

0.0%

0

L6

8287

H5

7143

L5

4797

Other values (463)

25499

更详细的信息

Toggle details

reclat

Numeric

Distinct count

Unique (%)

Missing (%)

Missing (n)

Infinite (%)

Infinite (n)

12739

27.9%

16.0%

7315

0.0%

0

Mean

Minimum

Maximum

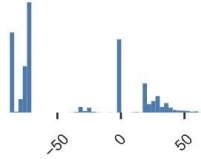
Zeros (%)

-39.10709514

-87.36667

81.16667

14.1%



Toggle details

强相关的会忽略并且提示

reclat_city

Highly correlated

This variable is highly correlated with reclat and should be ignored for analysis

Correlation

0.9942518712

可视化展示

reclong

Numeric

Distinct count

Unique (%)

Missing (%)

Missing (n)

Infinite (%)

Infinite (n)

14641

32.0%

16.0%

7315

0.0%

0

Mean

Minimum

Maximum

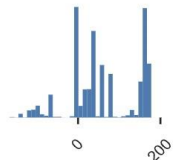
Zeros (%)

61.05259359

-165.43333


354.47333

13.6%



Pandas profiling

相关性



First rows

	boolean	fall	GeoLocation	id	mass_(g)	mixed	name	nametype	recclass	reclat
0	True	Fell	(50.775, 6.08333)	1	21.0	1	Aachen	Valid	L5	50.77500
1	False	Fell	(56.18333, 10.23333)	2	720.0	A	Aarhus	Valid	H6	56.18333
2	False	Fell	(54.21667, -113.0)	6	107000.0	1	Abee	Valid	EH4	54.21667
3	True	Fell	(16.88333, -99.9)	10	1914.0	A	Acapulco	Valid	Acapulcoite	16.88333
4	False	Fell	(-33.16667, -64.95)	370	780.0	1	Achiras	Valid	L6	-33.16667
5	False	Fell	(32.1, 71.8)	379	4239.0	A	Adhi Kot	Valid	EH4	32.10000
6	False	Fell	(44.83333, 95.16667)	390	910.0	A	Adzhi-Bogdo (stone)	Valid	LL3-6	44.83333
7	False	Fell	(44.21667, 0.61667)	392	30000.0	A	Agen	Valid	H5	44.21667
8	False	Fell	(-31.6, -65.23333)	398	1620.0	1	Aguada	Valid	L6	-31.60000
9	False	Fell	(-30.86667, -64.55)	417	1440.0	A	Aguila Blanca	Valid	L	-30.86667

前十行数据




2 Pandas under the hood



pandas

$$y_{it} = \beta' x_{it} + \mu_i + \epsilon_{it}$$



Python vs NDArray

- Python list

size

```
In [18]: import sys  
  
In [19]: lst = list(range(1000))  
  
In [20]: sys.getsizeof(lst)  
Out[20]: 9112
```

speed

```
In [30]: lst = list(range(100000))  
  
In [31]: %timeit sum((l + 1) for l in lst)  
7.9 ms ± 128 µs per loop
```

- Numpy NDArray

size

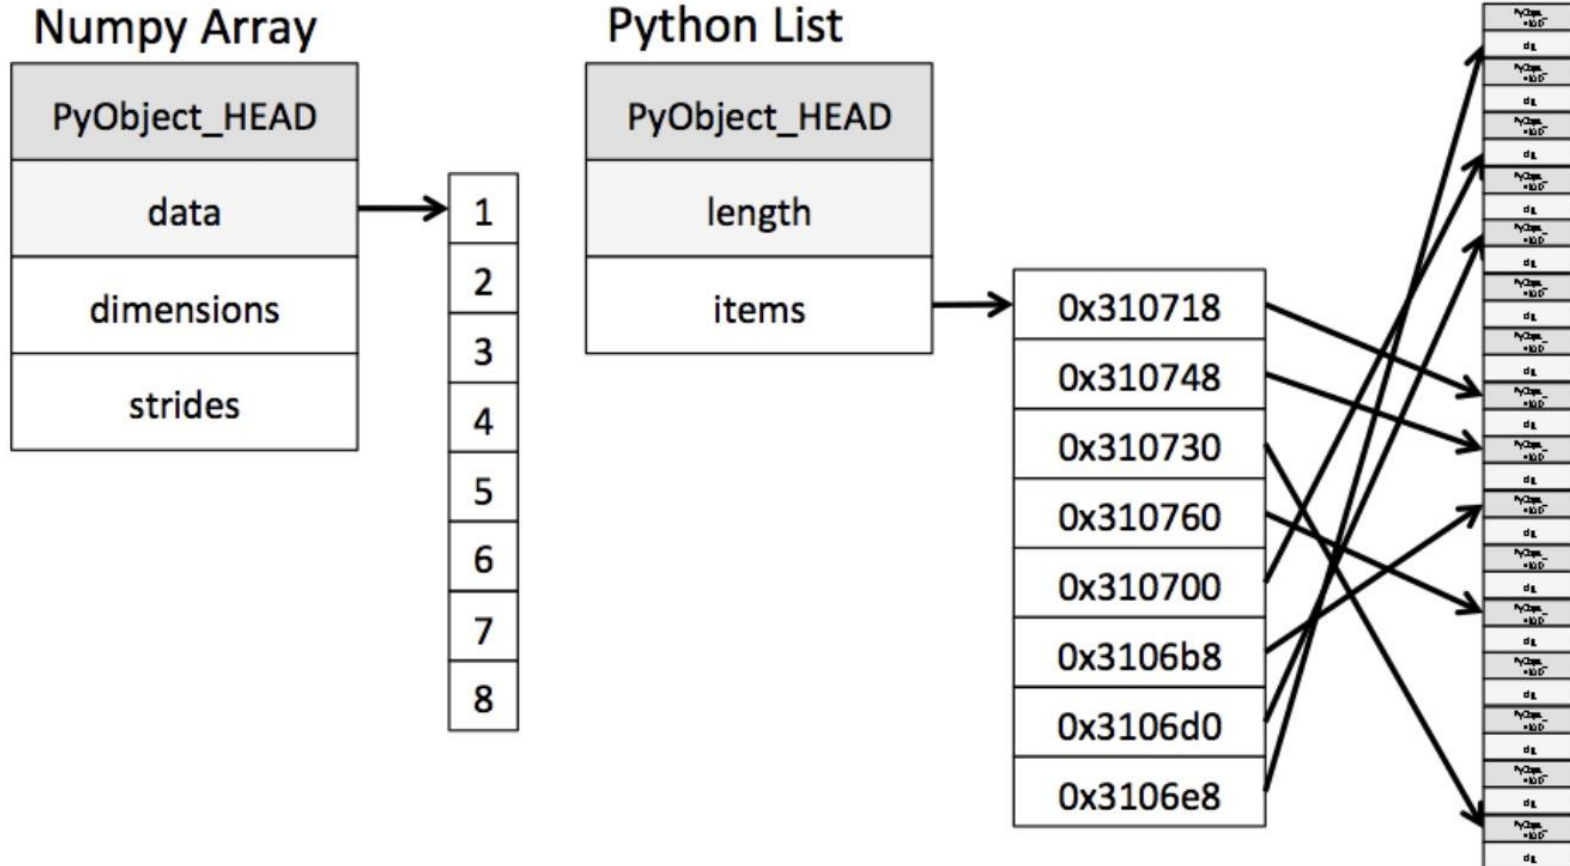
```
In [24]: import numpy as np  
  
In [25]: arr = np.arange(1000,  
dtype=np.int32)  
  
In [26]: sys.getsizeof(arr)  
Out[26]: 4096
```

speed

```
In [32]: arr = np.arange(100000,  
dtype=np.int32)  
  
In [33]: %timeit (arr + 1).sum()  
121 µs ± 2.32 µs per loop
```



Python vs NDAarray



Why python is slow





根据不同的数据类型分成不同的块（讲解一下 bytes）

DataFrame

	date	number_of_game	day_of_week	v_name	v_league	v_game_number	h_name	h_league	h_game_number	v_score	h_score	length_outs
0	01871054	0	Thu	CL1	na	1	FW1	na	1	0	2	54.0
1	18710505	0	Fri	BS1	na	1	WS3	na	1	20	18	54.0
2	18710506	0	Sat	CL1	na	2	RC1	na	1	12	4	54.0

IntBlock

	0	1	2	3	4	5
0	01871054	0	1	1	0	2
1	18710505	0	1	1	20	18
2	18710506	0	2	1	12	4

ObjectBlock

	0	1	2	3	4
0	Thu	CL1	na	FW1	na
1	Fri	BS1	na	WS3	na
2	Sat	CL1	na	RC1	na

FloatBlock

	0
0	54.0
1	54.0
2	54.0

<https://www.dataquest.io/blog/pandas-big-data/>



BlockManager



	name	id	score
0	wang	0	90.0
1	zhao	1	99.9
2	he	2	60.0

```
In [6]: df._data.blocks
Out[6]:
(FloatBlock: slice(2, 3, 1), 1 x 3, dtype: float64,
 IntBlock: slice(1, 2, 1), 1 x 3, dtype: int64,
 ObjectBlock: slice(0, 1, 1), 1 x 3, dtype: object)
```

```
In [7]: type(df._data.blocks[1].values)
Out[7]: numpy.ndarray
```

```
In [8]: df._data.blocks[1].values
Out[8]: array([[0, 1, 2]])
```





3 更高效的 Pandas

任务：收集了县级城市的经纬度，计算他们与杭州的距离。

```
In [3]: df
```

```
Out[3]:
```

	行政代码	地名	longitude	latitude
0	130426	河北省涉县	113.742914	36.598105
1	130930	河北省孟村回族自治县	117.159538	38.091265
2	140728	山西省平遥县	112.265493	37.148090
3	150424	内蒙古林西县	118.110216	43.771462
4	210202	辽宁省大连市中山区	121.677966	38.900436
...
3516	621201	甘肃省陇南市市辖区	104.934573	33.394480
3517	621202	甘肃省陇南市武都区	105.134553	33.293917
3518	621221	甘肃省成县	105.688289	33.747297
3519	621222	甘肃省文县	104.784206	32.947265
3520	621223	甘肃省宕昌县	104.452827	34.013489

```
[3521 rows x 4 columns]
```





haversine 公式:

```
def haversine(lat1, lon1, lat2, lon2):  
    # lat1, lon1 为位置 1 的经纬度坐标  
    # lat2, lon2 为位置 2 的经纬度坐标  
    import numpy as np  
  
    dlon = np.radians(lon2 - lon1)  
    dlat = np.radians(lat2 - lat1)  
    a = np.sin(dlat / 2) ** 2 + np.cos(np.radians(lat1)) * \  
        np.cos(np.radians(lat2)) * np.sin(dlon / 2) ** 2  
    c = 2 * np.arcsin(np.sqrt(a))  
    r = 6371 # 地球平均半径, 单位为公里  
    return c * r
```

```
# 杭州的经纬度为 (120.2E ,30.3N), 计算距离调用如下  
distance_hangzhou = haversine(lat, lon, 30.3, 120.2)
```



从循环开始说起

```
In [8]: %%timeit
...:
...: haversine_distances = []
...: for _, row in df.iterrows():
...:     distance = haversine(row['latitude'], row['longitude'], 30.3, 120.2)
...:     haversine_distances.append(distance)
...:
```

471 ms \pm 24.8 ms per loop (mean \pm std. dev. of 7 runs, 1 loop each)




Tip1: 避免循环

Apply

```
In [12]: %%timeit
...:
...: haversine_distances = df.apply(lambda row: haversine(row['latitude'],
...:                                                     row['longitude'], 30.3, 120.2), axis=1)
...:
```

146 ms \pm 4.21 ms per loop (mean \pm std. dev. of 7 runs, 10 loops each)





加速你的 Pandas

```
In [12]: %%timeit
...:
...: haversine_distances = df.apply(lambda row: haversine(row['latitude'],
...:                                                     row['longitude'], 30.3, 120.2), axis=1)
...:
146 ms ± 4.21 ms per loop (mean ± std. dev. of 7 runs, 10 loops each)
```



Tip2: 直接对列操作

精通面向数组的编程和思维方式是成为
Python 科学计算牛人的一大关键步骤
—— Wes McKinney

```
In [19]: %%timeit
...:
...: df['dis'] = haversine(df['latitude'], df['longitude'], 30.3, 120.2)
...:
...:
```

2.62 ms ± 286 μs per loop (mean ± std. dev. of 7 runs, 100 loops each)



加速你的 Pandas

还不够快?

```
In [20]: %%timeit
...:
...: df['dis'] = haversine(df['latitude'].to_numpy(),
...:                      df['longitude'].to_numpy(), 30.3, 120.2)
...:
...:
```

376 μ s \pm 7.55 μ s per loop (mean \pm std. dev. of 7 runs, 1000 loops each)

471 ms \pm 24.8 ms per loop

1000 倍以上的
提升



Tip3: 数组计算直接使用 numpy





Categorical data

- 颜色: Red, Black, White, Green, ...
- 性别: Male, Female
- 星期日期: Monday, Tuesday, Wednesday, ...
-

绝大部分分类数据是字符串

Object 类型意味着大内存与低效

DataFrame

	date	number_of_game	day_of_week	v_name	v_league	v_game_number	h_name	h_league	h_game_number	v_score	h_score	length_outs
0	01871054	0	Thu	CL1	na	1	FW1	na	1	0	2	54.0
1	18710505	0	Fri	BS1	na	1	WS3	na	1	20	18	54.0
2	18710506	0	Sat	CL1	na	2	RC1	na	1	12	4	54.0

IntBlock

	0	1	2	3	4	5
0	01871054	0	1	1	0	2
1	18710505	0	1	1	20	18
2	18710506	0	2	1	12	4

ObjectBlock

	0	1	2	3	4
0	Thu	CL1	na	FW1	na
1	Fri	BS1	na	WS3	na
2	Sat	CL1	na	RC1	na

FloatBlock

	0
0	54.0
1	54.0
2	54.0





```
In [28]: df['省份']  
Out[28]:  
0      河北  
1      河北  
2      山西  
...  
3518   甘肃  
3519   甘肃  
3520   甘肃  
Name: 省份, Length: 3521, dtype: object
```

用 category 表示省份

```
In [29]: df['省份'] = df['省份'].astype('category')  
  
In [30]: df['省份']  
Out[30]:  
0      河北  
1      河北  
2      山西  
..  
3518   甘肃  
3519   甘肃  
3520   甘肃  
Name: 省份, Length: 3521, dtype: category  
Categories (33, object): [上海, 云南, 内蒙, 北京, ..., 陕西, 青海, 香港特别行政区, 黑龙江]
```



```
In [29]: series_as_category = df['省份']
Out[29]:
0      河北
1      河北
2      山西
..
3518   甘肃
3519   甘肃
3520   甘肃
Name: 省份, Length: 3521, dtype: category
Categories (33, object): [上海, 云南, 内蒙, 北京, ..., 陕西, 青海, 香港特别行政区, 黑龙江]
```

实际上已经映射成整型 (int8)

```
In [66]: series_as_category.cat.codes
Out[66]:
0      16
1      16
2      10
..
3518   23
3519   23
3520   23
Length: 3521, dtype: int8
```





Tip4: 使用 category 类型

- 内存节约 (~5.4 x)

```
In [49]: series = df['省份']
```

```
In [50]: series.memory_usage()
```

```
Out[50]: 28296
```

```
In [51]: series_as_category = series.astype('category')
```

```
In [52]: series_as_category.memory_usage()
```

```
Out[52]: 5193
```

- 性能 (~1.8 x)，数据量很大时性能会有大幅提升

```
In [58]: %timeit df.groupby('省份')['dis'].max()  
913 µs ± 18.3 µs per loop (mean ± std. dev. of 7 runs, 1000 loops  
each)
```

```
In [59]: %timeit df_category.groupby('省份')['dis'].max()  
510 µs ± 15.8 µs per loop (mean ± std. dev. of 7 runs, 1000 loops  
each)
```



当数据规模来到了 100G，所有的优化都是白费，

求助 SQL? Spark?

不，我们需要并行的、分布式的 Pandas!





4 更大规模的数据分析



Mars: 一个统一的分布式计算框架，将 Python 科学技术栈应用到大数据中

- Numpy
- Pandas
- Scikit-learn

```
In [1]: import pandas as pd
```

```
In [2]: import numpy as np
```

```
In [3]: df =  
pd.DataFrame(np.random.rand(10, 10))
```

```
In [4]: df.sum()
```

```
Out[4]:
```

```
0    4.639573
```

```
1    5.552392
```

```
2    4.974636
```

```
3    3.709493
```

```
4    3.886881
```

```
5    4.640418
```

```
6    5.660901
```

```
7    5.748306
```

```
8    4.227453
```

```
9    5.897225
```

```
dtype: float64
```

相同的 API



```
In [1]: import mars.tensor as mt
```

```
In [2]: import mars.dataframe as md
```

```
In [3]: df =  
md.DataFrame(mt.random.rand(10, 10))
```

```
In [4]: df.sum().execute()
```

```
Out[4]:
```

```
0    4.657526
```

```
1    6.373730
```

```
2    5.526396
```

```
3    3.830616
```

```
4    5.070046
```

```
5    4.986147
```

```
6    5.586429
```

```
7    4.408413
```

```
8    4.461475
```

```
9    4.800761
```

```
dtype: float64
```



```
In [5]: df1 = pd.DataFrame(np.random.rand(10, 3))
```

```
In [6]: df2 = pd.DataFrame(np.random.rand(10, 3),  
index=[4, 1, 3, 2, 10, 5, 9, 8, 6, 7])
```

```
In [7]: df1 + df2
```

```
Out[7]:
```

	0	1	2
0	NaN	NaN	NaN
1	0.523616	0.877549	1.129148
2	1.420484	1.176361	1.504610
3	1.137251	0.582319	0.545606
4	1.433731	1.271197	1.069266
5	1.573455	1.078515	1.509255
6	0.916412	0.511708	0.741676
7	1.284630	0.579828	1.366145
8	1.283414	0.677225	0.136293
9	0.453015	1.801961	1.358312
10	NaN	NaN	NaN

```
In [5]: df1 = md.DataFrame(mt.random.rand(10, 3))
```

```
In [6]: df2 = md.DataFrame(mt.random.rand(10, 3),  
index=[4, 1, 3, 2, 10, 5, 9, 8, 6, 7])
```

```
In [7]: (df1 + df2).execute()
```

```
Out[7]:
```

	0	1	2
0	NaN	NaN	NaN
1	0.342513	0.458722	1.106504
2	0.776831	0.549000	1.187729
3	0.483544	0.564384	0.245915
4	0.993470	0.102466	1.615469
5	0.744108	1.019650	1.522074
6	0.671610	0.820533	1.150278
7	0.826367	0.872842	0.877645
8	0.503570	1.328384	0.987532
9	1.654519	1.026768	1.456997
10	NaN	NaN	NaN

一致的行为，自动索引对齐





```
In [10]: %%time
...: df = pd.DataFrame(np.random.rand(100000000, 4))
...: df.sum()
...:
...:
CPU times: user 12 s, sys: 3.6 s, total: 15.6 s
Wall time: 12.8 s
Out[10]:
0    4.999985e+07
1    4.999891e+07
2    4.999715e+07
3    5.000027e+07
dtype: float64
```

运行时间: 12.8s

内存峰值: 3430.29M

单机上, 自动利用多核,
加速计算

```
In [4]: %%time
...: df = md.DataFrame(mt.random.rand(100000000, 4))
...: df.sum().execute()
...:
...:
CPU times: user 19.3 s, sys: 5.14 s, total: 24.4 s
Wall time: 4.27 s
Out[4]:
0    4.999827e+07
1    5.000004e+07
2    4.999853e+07
3    4.999748e+07
dtype: float64
```

运行时间: 4.27s

内存峰值: 2007.92M



Mars

分布式

部署 Mars 集群，通过
RESTful API 提交任务

```
In [5]: from mars.session import new_session
```

```
In [6]: sess = new_session('http://127.0.0.1:40002').as_default()
```

```
In [7]: df = md.DataFrame(mt.random.rand(100000000, 4))
```

```
In [8]: df.sum().execute()
```





```
import numpy as np
```

```
def haversine(lat1, lon1, lat2, lon2):  
    # lat1, lon1 为位置 1 的经纬度坐标  
    # lat2, lon2 为位置 2 的经纬度坐标  
  
    dlon = np.radians(lon2 - lon1)  
    dlat = np.radians(lat2 - lat1)  
    a = np.sin(dlat / 2) ** 2 +  
np.cos(np.radians(lat1)) * \  
        np.cos(np.radians(lat2)) * np.sin(dlon / 2)  
    ** 2  
    c = 2 * np.arcsin(np.sqrt(a))  
    r = 6371 # 地球平均半径, 单位为公里  
    return c * r
```

```
%time haversine(latitude, longitude, 30.3, 120.2)
```

```
CPU times: user 11.5 s, sys: 9.6 s, total: 21 s
```

```
Wall time: 22 s
```

```
import mars.tensor as np
```

```
def haversine(lat1, lon1, lat2, lon2):  
    # lat1, lon1 为位置 1 的经纬度坐标  
    # lat2, lon2 为位置 2 的经纬度坐标  
  
    dlon = np.radians(lon2 - lon1)  
    dlat = np.radians(lat2 - lat1)  
    a = np.sin(dlat / 2) ** 2 +  
np.cos(np.radians(lat1)) * \  
        np.cos(np.radians(lat2)) * np.sin(dlon / 2)  
    ** 2  
    c = 2 * np.arcsin(np.sqrt(a))  
    r = 6371 # 地球平均半径, 单位为公里  
    return (c * r).execute()
```

```
%time haversine(latitude, longitude, 30.3, 120.2)
```

```
CPU times: user 20.4 s, sys: 19.2 s, total: 39.5 s
```

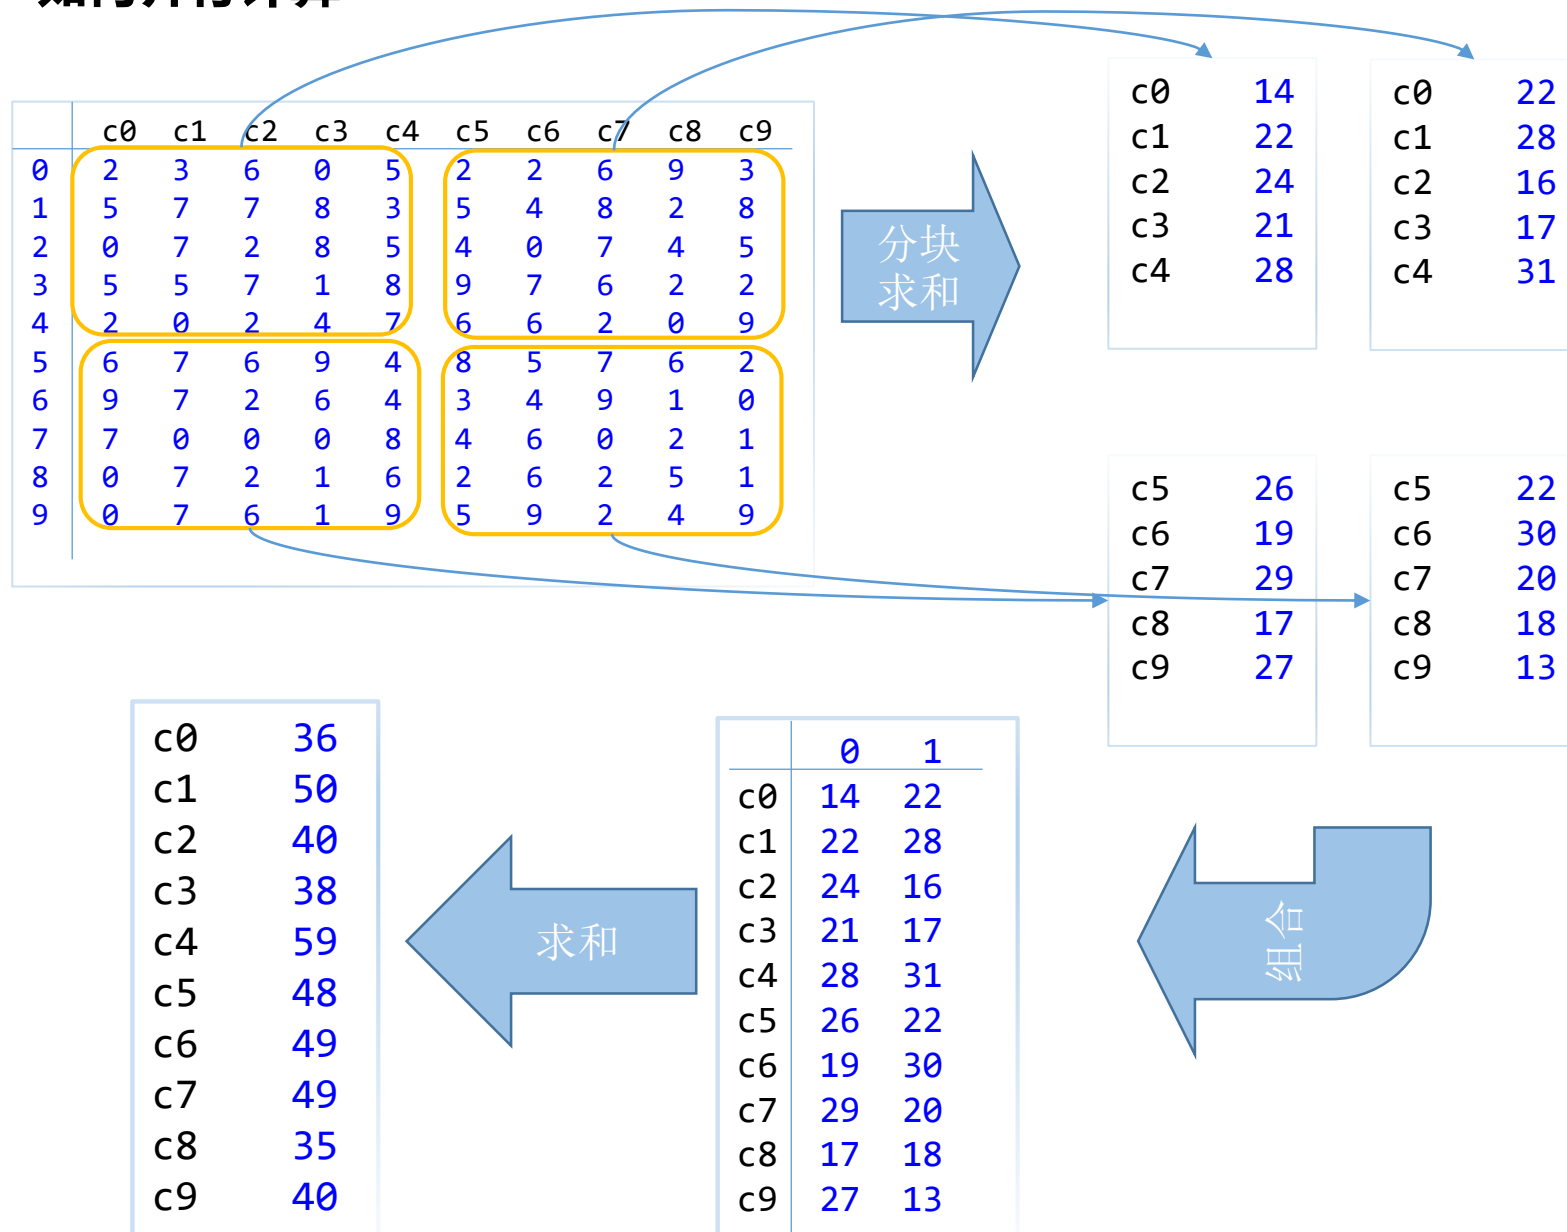
```
Wall time: 13.2 s
```



Tip5: 使用 Mars



如何并行计算

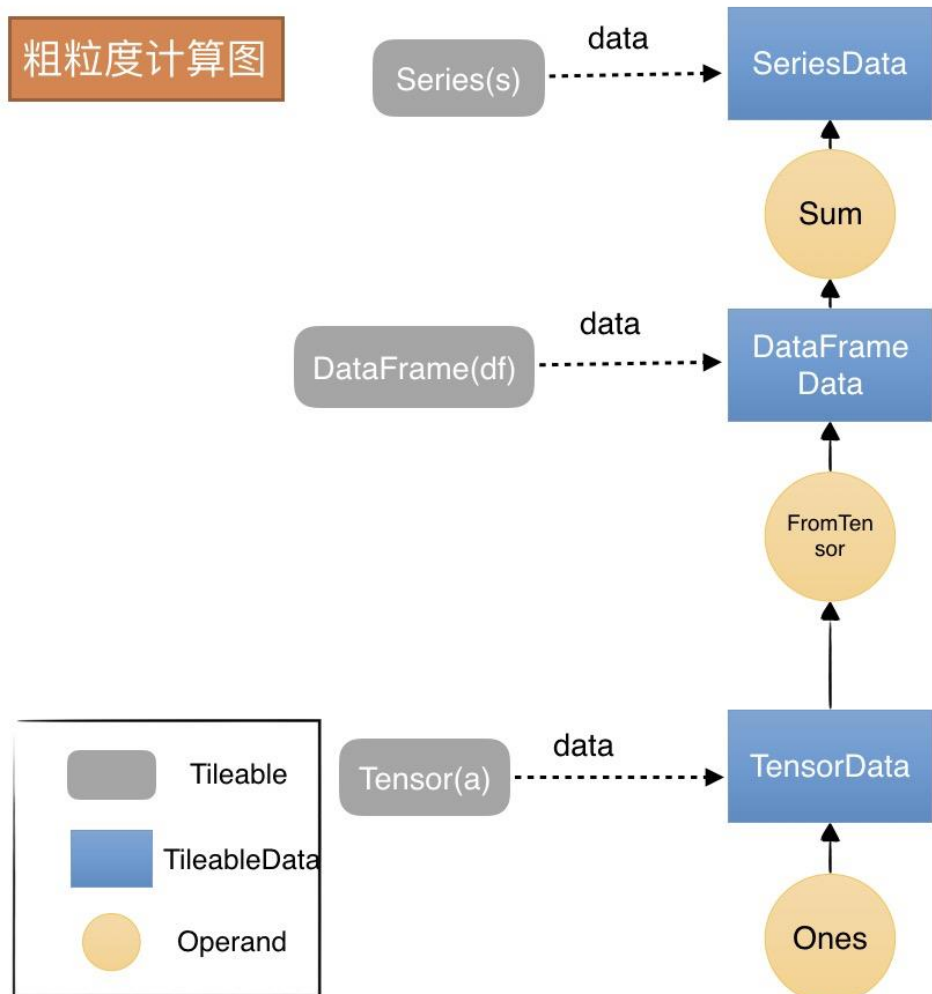


如何并行计算

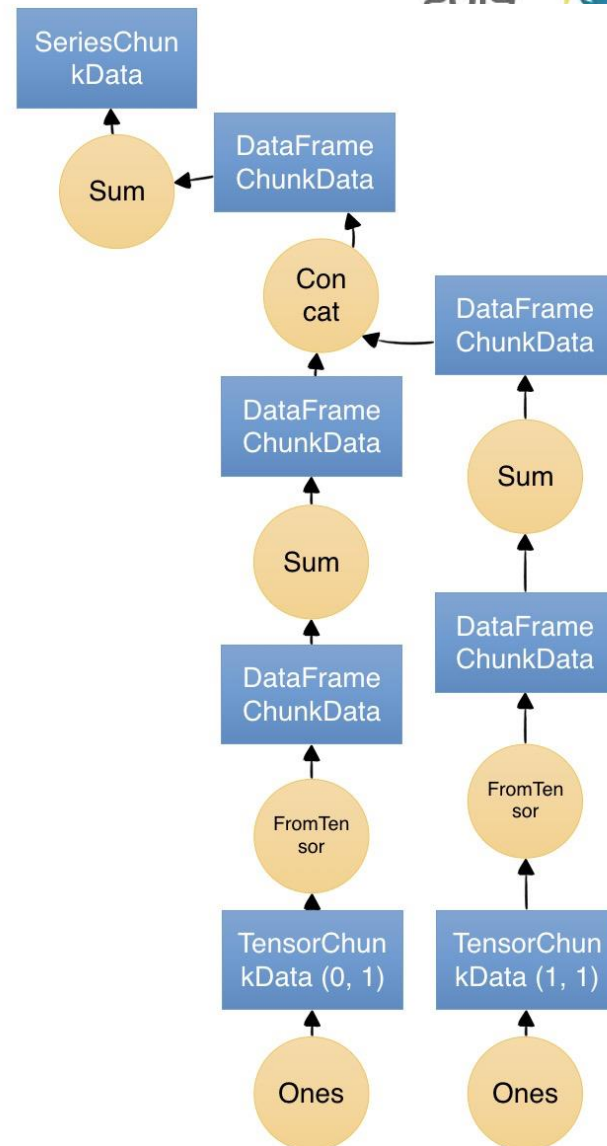
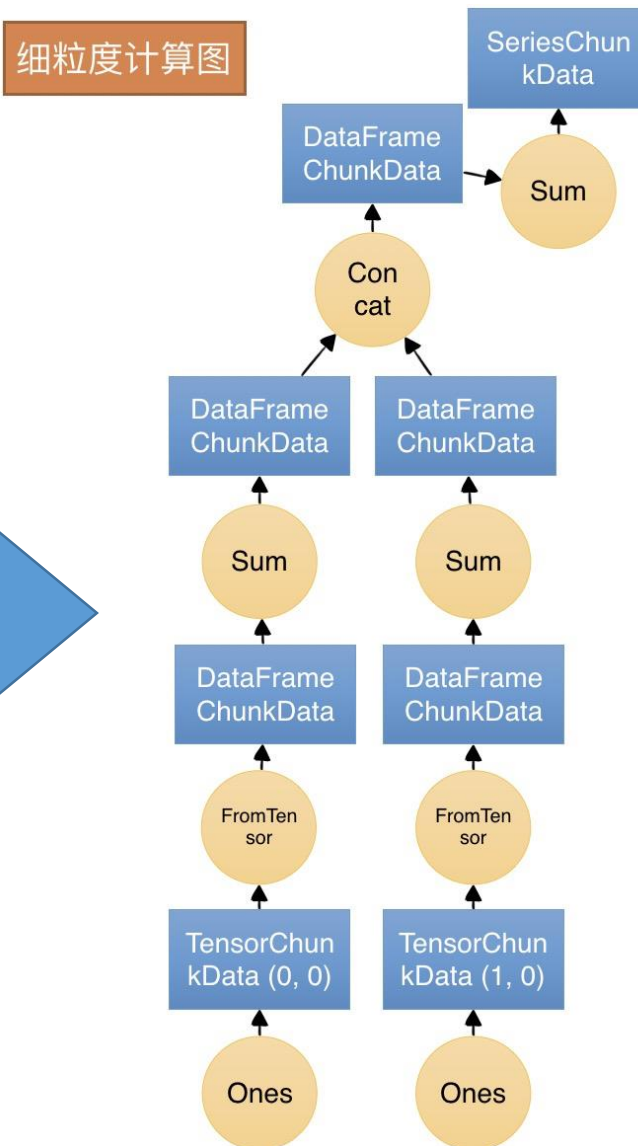


分而治之

```
In [1]: import mars.tensor as mt
In [2]: import mars.dataframe as md
In [3]: a = mt.ones((10, 10),
chunk_size=5)
In [4]: df = md.DataFrame(a)
In [5]: s = df.sum()
In [6]: s.execute()
Out[6]:
0    10.0
1    10.0
2    10.0
3    10.0
4    10.0
5    10.0
6    10.0
7    10.0
8    10.0
9    10.0
dtype: float64
```

粗粒度计算图



分而治之





Mars DataFrame

- 实现的接口
 - 创建 DataFrame: DataFrame、from_records
 - Basic arithmetic: 基本算数运算
 - Math: 数学运算
 - Indexing: 索引
 - iloc
 - 列选择
 - set_Index
 - Reduction: 聚合
 - Groupby: 分组聚合
 - merge/join
- Roadmap: <https://github.com/mars-project/mars/issues/495>



Mars 现状及方向

- `pip install pymars`
- Mars 开源地址: <https://github.com/mars-project/mars>
- Mars 文档: https://docs.mars-project.io/zh_CN/latest/
- 双版本发布
- 方向:
 - 社区是重点
 - 技术:
 - Roadmap 和 Enhancement Proposal:
<https://github.com/mars-project/mars/issues/537>
 - 丰富 DataFrame、learn 和 Tensor 的接口
 - 更好的算子融合
 - Mars actors 层优化, 支持更高效的执行和网络效率
 - 支持更多调度





THANK YOU



Pole

浙江 杭州

